

Streamline — Project Status Report

Prepared for Angel · Delivery lead: Claude Code · Report date: 2026-08-23 · First issue of this report

1. Executive Summary

Overall status: **AMBER**

This is the first status report issued against Streamline. There is no prior period to compare against, so this report establishes the baseline. All 27 decision records on file, the full data-extraction rebuild, and the entire design system were produced in a single extended working session rather than across a calendar week; the commit history and the working tree both reflect that, and the discrepancy is called out explicitly in section 2 rather than smoothed over.

The rating is amber rather than green for two reasons, neither of which is a failure of delivery. First, a materially wrong figure was found and corrected today on the one filer that was rendering before this period began (section 2, and decision record 0016 for the earlier instance of the same failure mode). The correction is verified and shipped, but a product whose stated guarantee is that no displayed figure is invented had, for a period, displayed one that was wrong by \$454 million. Second, the product's core visual metaphor has hit a case it does not yet have a defined answer for — a business segment that loses money inside a profitable company — and that case is common in the technology sector, not an edge case. A proposal is in front of the product owner; nothing ships against it until a decision is made.

It is not red because both findings were caught by the process built to catch them, before either reached a user, and both are fully diagnosed with a fix in hand or a decision teed up. It is not green because the coverage gate has never completed a full pass and two of six items on the decision queue have sat open since this report's cutoff.

2. Progress This Period

The product went from rendering one company correctly (Microsoft) to rendering six companies correctly (Microsoft, Apple, NVIDIA, HP, Autodesk, Meta), against a background corpus of 19 real filings pulled from the SEC and used as the standing test set going forward.

2.1 Extraction pipeline

The team that turns a raw SEC filing into the numbers behind the picture found and closed five defects, all discovered by testing real companies rather than the one company the product had been built and demonstrated against.

Defect	Effect before fix	Status
Autodesk's reported operating income was read incorrectly	The app displayed \$1,124M against a filed figure of \$1,578M — a \$454M error, live on the only filer besides Microsoft that rendered at all	Fixed and verified against the filing
Income tax was drawn as a business cost	A river's width included tax, which is a company-wide item, not a segment cost — same root cause as the figure above	Fixed
A recognized industry-standard footnote (a company disclosing one figure to two decimal precisions) was treated as a contradiction	IBM's entire annual report was rejected outright over a rounding difference of roughly one part in 10 ¹⁷	Fixed

Defect	Effect before fix	Status
An allowed disclosure pattern could never actually be recognized by the code (a one-line defect)	Apple was rejected; the rule meant to allow its disclosure style could never fire	Fixed — Apple now renders
A company's segment total was only read from the first matching entry in the filing	HP rendered correctly by chance, because of the order data happened to appear in its filing; a different filing order would have produced a wrong figure with no warning	Fixed — now reads and cross-checks every matching entry

A further rule change, completed as this report was being finalized, removed a second class of risk before it could produce a wrong figure: a company whose disclosed costs do not add up to its reported profit will now be refused rather than rendered with a warning. Every company currently working was re-checked against this new rule before it shipped, and none of the sixteen segments across the six companies was affected — the rule is a trap for the next company, not a change to today's picture.

2.2 Coverage rule

The product owner approved widening the industry-coverage rule to include semiconductor companies (adds NVIDIA, AMD, Intel, and similar names). A related but different code — the one covering ride-share and other miscellaneous business services, which would have added Uber — was considered and explicitly declined, on the grounds that the code in question is SEC's general catch-all for businesses that do not fit elsewhere, not a technology classification. The distinction is recorded so it is not revisited as an oversight.

2.3 Visual design

The application moved from unstyled browser defaults to a complete, deliberate visual system: a typeface chosen and licensed for the product, a color palette derived from measurements of real water photography rather than picked by eye, and a full set of interface components. This was built from a written design brief (five reference sites were measured, not guessed at) that the product owner approved before any of it was built.

2.4 Team structure

The nine working roles on this project were renamed to plain job titles (for example, the person who used to be called “Ledger” is now the Financial Data Analyst), and two new roles were added: a DevOps Engineer, for how the product is built and deployed, and a Reliability Engineer, for knowing when something is wrong in production before a customer reports it. Neither new role has started work yet.

2.5 Automated checks

Check	Result
Automated test suite	1,000 of 1,000 checks passing, across 80 test files
Type-correctness check	Clean, zero errors
Animation smoothness (99th percentile of a large sample of frames, not an average)	17.6 milliseconds against a 16.7 millisecond budget for 60-frames-per-second — within the built-in safety margin; the product is designed to step down to a slower, still-smooth rate rather than stutter
Companies verified to render correctly against the real SEC filing	6 of 19 in the test corpus (up from 1 of 19 at the start of this period)

Note on the Meta Platforms case specifically, because it was described inaccurately in an earlier verbal update today: Meta does **not** currently render. Its Reality Labs division spent more than nine times what it earned, which would require drawing a business segment as a negative-width shape — something the product has no defined way to draw honestly.

Rather than fake it, the product currently shows a plain, correct explanation of why it will not draw that company yet. That is the right behavior today. The proposal to give this case a real, honest drawing is in section 4.

3. Decision Register

27 decision records exist on file, all produced during this period; there is no earlier decision history to report against. The table below lists every decision resolved this period with the product owner's ruling; the open-items table follows.

3.1 Resolved this period

#	Decision	Ruling
—	How should the lake (cash kept) relate to the river feeding it?	Keep them visually separate for now; do not imply a connection the numbers cannot support. A fuller fix is the long-term goal but is bigger than this phase of work.
—	A wrong number was found live in production — how strict should the fix be?	The original filed number wins, always, even if a summary elsewhere says otherwise.
—	Should the web address include a specific reporting period (quarter, year)?	Not yet — deliberately deferred; a test now fails on purpose if anyone adds one prematurely.
—	What happens visually when a company's segment completely runs out of money?	The picture stops cleanly at that point rather than continuing on with fabricated numbers. A placeholder label showing \$0 stays visible on purpose, as an honest flag that real wording is still needed.
—	Approve a new testing tool (Playwright) to measure real animation speed?	Approved. The alternative (a fake in-memory browser) cannot produce trustworthy speed numbers.
—	Which typeface for the product?	IBM Plex (free, professionally licensed, self-hosted — no ongoing cost or vendor dependency).
—	Should industry coverage include semiconductor companies?	Yes — approved. Should it include the general miscellaneous-business-services code that would sweep in Uber? No — explicitly declined as too broad.
—	How should a business segment with weak internal cost math be handled?	Refuse to draw it rather than draw it with only a warning. Verified this affects zero of the six companies currently working.

3.2 Still open, awaiting the product owner

Decision	What it blocks	Age
How to draw a business segment that loses more money than it earns, inside an otherwise profitable company (the Meta Platforms / Reality Labs case)	Meta Platforms cannot render at all until this is answered; a written proposal with three concrete options is ready for review — see section 4	New this period
When a company's tax bill and other non-operating items are combined into one number in the picture, should they be shown separately instead?	No company is currently blocked by this, but Microsoft's own numbers already show why it matters: the combined number understates the actual tax bill by roughly \$10.7 billion	1 period
What should the app show before a user has picked a company (the very first screen)?	Depends on two other open decisions below (default industry scope and default reporting period); nothing can be finalized until those are answered	1 period
Is the current industry-coverage rule (technology sector by SIC code) the right one long-term?	Determines how many real companies are eligible at all; also affects the wording shown to a user searching for a company outside coverage	1 period

Decision	What it blocks	Age
Which reporting period should load by default when a user opens the app — most recent year, most recent quarter, or trailing twelve months?	Blocks the first-screen design and the web address format	1 period
Should an amended filing (a company's own correction to a prior report) be flagged to the user, or silently treated as if it were the original?	No company in the current test set has an amended filing, so nothing is broken today — but the current code silently ignores the correction if one exists	1 period

None of the six items above has appeared in a prior report, so none has crossed the three-report escalation threshold this routine watches for. That threshold starts counting from this report.

4. Risk Register

Risk	Likelihood	Impact	What is being done
A future correctness fix to how figures are calculated is silently ignored because the app is still serving an old, cached answer	Confirmed — happened today	High — directly undermines the product's core promise that every number is real	Diagnosed and reported to the team that owns it; a fix is scoped but not yet built. No existing safeguard catches this automatically today.
The visual metaphor (rivers and lakes) does not have a defined, honest way to draw a loss-making segment inside a profitable company	High — common in the technology sector, not rare	High — blocks an unknown but likely large share of real companies from ever rendering	A written proposal with three options and a recommendation is ready for the product owner — see section 4 below the register
No formal adversarial testing team has completed a full pass; today's testing folder is still nearly empty of independently-authored tests	Confirmed	Medium — most defects are still being found by the same people who wrote the code being tested	First outside test pass is in progress as of this report; initial results (six real companies verified) are already folded into section 2 above
No automatic, unattended testing pipeline exists yet — nothing runs the checks automatically when code changes	Confirmed	Medium — a broken change could reach the working copy without being caught immediately	Deliberately not started this period; the product owner asked to handle version control personally first, and pipeline work is sequenced to start once that is done
Two hundred and more changed files exist only on this machine, with no backup or history saved	Confirmed	High if this machine is lost or damaged — the day's entire body of work would be unrecoverable	Flagged to the product owner; owner has stated this is being handled personally
A missing folder that the app depends on causes a confusing technical error instead of a plain one	Low — requires an unusual setup mistake	Low	Found and reported; not yet fixed

5. Blockers and Decisions Required From the Product Owner

Ordered by how much they are currently holding up. Each has a recommended next step; none requires an immediate response, but the first two are worth prioritizing.

Decision needed	Recommended next step	Suggested by
How should a money-losing business segment be drawn inside an otherwise profitable company?	Review the three-option proposal (section 4 below) and choose one, or ask for a fourth option	This week
Save the day's work to version control (or confirm someone else will)	A single save-point now would prevent total loss if this machine has a problem	This week
Should the tax bill and other non-operating items be shown as one combined number in the picture, or broken out separately?	Low urgency — no company is blocked by this today, but it is worth deciding before more companies are added	No firm deadline
Should companies outside today's technology-sector rule eventually be supported, and if so, which ones?	Not urgent; affects long-term scope more than this week's work	No firm deadline

Proposal on file: the loss-making-segment question

Three options were written up by the visualization team, each stated in terms of what a viewer would correctly understand and what they might misread. The short version:

Option A (recommended): Close that segment's river cleanly at the point its costs run out, and show the shortfall as a separate, clearly labeled dip in the combined flow right after the rivers merge. Every dollar is still accounted for and visible; the loss is drawn at a size a viewer can actually see, rather than being invisible.

Option B: Divert the shortfall into a second, separate reservoir shape, similar to how an entire loss-making company is shown today. Rejected as the primary choice because it would put two different kinds of “water pool” on screen side by side, at sizes that are not really comparable, inviting a viewer to compare two different measurements as if they were the same thing.

Option C: Do nothing further — keep declining to draw companies with this pattern, as it does today. Costs nothing extra, but rules out a meaningful share of real technology companies from ever appearing in the product.

The visualization team's recommendation is Option A. Nothing will be built against any option until the product owner chooses.

6. Recommendation for the Coming Period

- Decide the loss-making-segment question (section 4) — it is the single item most likely to unblock the largest number of additional real companies.
- Save today's work to version control, even as a single rough save-point, before any further work continues on this machine.
- Let the independent testing pass currently underway finish and report its full findings before adding more companies to the working set — it is the only check in the process that does not report to the same people who built the thing being checked.
- Fix the caching problem that let a stale, wrong answer keep being served after the correct fix had already landed — this is the highest-impact defect on file precisely because it could hide a future correct fix the same way.
- Hold off on starting the automatic deployment pipeline until the product owner has finished the personal version-control work already in progress.

This report reflects only what could be read directly from the project's files, its test results, and its written decision records at the time it was produced. Where a record was missing or a claim could not be verified against a file, that gap is stated above rather than estimated.