

Streamline — Project Status Report

Prepared for Angel · Delivery lead: Claude Code · Report date: 2026-09-01 · Second issue · Period covered: 2026-08-23 12:49 to 2026-09-01

1. Executive Summary

AMBER

Status unchanged from the first issue, but the composition of the amber has changed substantially. One of the two original causes is closed; a new one has opened in its place.

The single most important change since the first issue is that the work is no longer at risk of total loss. The prior report asked for a save-point to version control as a this-week item. That is done and then some: pull request 5 was opened, reviewed and merged into **main** (merge commit 8234722), carrying 52 files and 2,047 added lines. The working branch now sits level with main at zero commits ahead. That risk is closed and does not carry forward.

Against that, the performance gate has moved in the wrong direction. At the last issue the position recorded in STATUS.md was that every absolute Invariant 4.1 gate passed with wide margin and only the stale baseline comparison tripped. On today's run, two of nine checks fail, and one of them is an absolute gate rather than a comparison. The margins are small and the evidence points at the measurement environment rather than the build, but the honest statement is that the project cannot currently produce a clean performance run, and a product whose fourth invariant is a frame budget should be able to.

Three qualifications belong beside that finding, and section 5 of this report does not soften any of them. First, the configuration that real filers actually hit — the capped load, at most eight lanes plus an aggregate under Invariant 3.7 — passed. Only the uncapped twelve-lane synthetic stress case failed, and it failed by 0.1 ms against a 25.0 ms limit, a margin of four parts in a thousand. Second, the regression comparison trips against a baseline that was itself re-recorded during this period on a machine that was not idle, so the reference number may be contaminated. Third, today's run was taken while this same machine was executing the 1,075-test suite and hosting three local servers. The measurement conditions were poor and that is a fault in the measurement, not an exoneration of the build.

The other original cause of the amber rating is unchanged. A money-losing business segment inside a profitable company still has no defined visual treatment, Meta still refuses to draw because of it, and the decision has now been outstanding across two consecutive reports. Under this report's own escalation rule a third consecutive miss is raised into this summary as a headline item; one more period without an answer reaches that threshold.

Delivery itself was strong. Three decision records were ruled and shipped, the visual direction question that had been open since 2026-08-23 is now half-answered and built, and the test suite grew from 1,032 to 1,075 with every check green alongside a clean typecheck and formatter. Nothing in this report should be read as a delivery problem. The rating is amber because verification is weaker than delivery, and because the product owner is now the critical path on three separate questions.

2. Progress This Period

Two commits landed in the reporting window: 4c22804 on 2026-08-23 at 15:12 and cd98ab5 on 2026-08-30 at 14:09. That count is not the measure of the period and is not offered as one. The second commit alone carries 52 changed files and represents three separate working sessions consolidated before merge. What follows is stated in verdicts and artifacts rather than in volume.

Shipped and verified

Item	Evidence	State
0035 — the lake is the company as a whole. The definition of the lake was naming its quantity rather than its meaning. Net earnings is now recorded as its magnitude, not its meaning. Invariants 1 and 3.3 amended with an amendment-log entry.	docs/decisions/0035 · commit 4c22804	Complete. No geometry, scale or pipeline change; binding on copy and the analyst panel, neither of which is built.
0037 — the water reads as water. The saturation ceiling set by 0022 was amended by ruling. The water ramp is now genuinely blue while each step keeps its measured photographic luminance; river, trunk and lake share one hue family and one memoized gradient fill.	docs/decisions/0037 · commit cd98ab5	Complete. Colour guards were rewritten to enforce the new rule rather than deleted: ceiling 65 per cent, one shared hue family 195–215°, no per-segment variation.
0038 — the world dressing. Every covered filer now renders inside a sunset sky band with hill silhouettes, dusk terrain, mist and a water rim glow, drawn around unchanged financial geometry.	docs/decisions/0038 · src/viz/render/world.ts · commit cd98ab5	Complete. Scenery is seeded by the company identifier alone; tests assert that doubling every filed figure moves no scenery pixel and that refusal states render with no world painted.
Figures made readable. Canvas figures now read compact — \$133.7B — with the exact filed value, \$133,749M, one hover away. The hover overlay gained rows and chrome; secondary prose moved off the canvas into a document margin panel.	src/viz/render/format.ts · commit cd98ab5	Complete. Recorded as a deliberate change to how test record 0001 condition C2 is satisfied, since the persistent figure is now the rounded one.
Version control. The prior report's this-week request.	PR 5 · merge commit 8234722	Complete and closed as a risk.

Verification state at the close of the period

Check	Result	Comparison to prior issue
Automated test suite	1,075 passed across 86 files; zero failures	Up from 1,032 across 83 files
Type checking	Clean across both project configurations	Unchanged
Formatter	Clean	Unchanged
Performance gate	7 of 9 passed; 2 failed	Regression, see section 5

Frame-time distribution, measured 2026-09-01

Percentiles are given in full rather than as averages. Three load scenarios were measured. The right-hand column is the renderer's own drawing cost per frame, which is the figure that would rise if the world dressing

were responsible for the misses.

Scenario	Locked	p50	p95	p99	Worst	n	Draw p99 / worst	Verdict
Capped reference load, 9 lanes	16.67	16.70	16.80	17.60	17.70	601	0.40 / 2.70	PASS
Uncapped stress, 12 lanes	16.67	16.70	16.80	25.00	25.10	600	0.40 / 0.40	FAIL
Pinned to the 30 fps floor	33.33	41.70	41.80	50.10	50.10	80	0.70 / 0.70	FLAGGED

Two readings deserve emphasis. The renderer's own drawing cost never exceeds 0.7 ms at the 99th percentile against budgets of 16.67 and 33.33 ms, so the work Streamline performs per frame is using roughly four per cent of the time available; the misses are in the time between frames, not in the drawing. Separately, the third scenario reports a median frame interval of 41.70 ms while the rate-lock check simultaneously reports a locked 30.00 Hz, which is 33.33 ms. Those two statements cannot both describe an idle machine. Taken with the draw costs, the most probable reading is that the host was contended throughout, which matches the conditions: the full test suite and three local servers were running during the measurement. Interaction responsiveness was unaffected, measuring 17.80 ms at the 95th percentile against a 100 ms standard.

Workstream position against the delivery plan

The workstream table in STATUS.md is unchanged from the first issue. Workstreams 0 through 6 — scaffold, ingestion, extraction, derivations, lake encoding, constriction rendering and the performance harness — remain complete. Workstream 7, fixtures and regression, remains not started, and the QA Engineer role has still not been invoked. DevOps and Reliability remain not started. No workstream advanced this period; the work that landed was ruled visual direction and correctness inside already-complete workstreams.

Work outside the governed repository

A public-facing marketing site was built during this period and is not part of the Streamline repository or its agent protocol. It comprises a scroll-driven hero film, six photographic sections, an in-page viewer that draws live filers from the running application, and a single combined local server that now serves the site, the built application and the SEC proxy from one address. Measured on the live local build: 399 KB initial transfer across 6 requests with a 613 ms load event, with the 6.1 MB hero film streaming behind a progress ring after the page is usable. It is reported here for completeness and because it carries a risk noted in section 4; it is not counted as workstream progress.

3. Decision Register

Thirty-seven decision records are on file, numbered to 0038. Three were added this period. One number, 0036, does not exist: two records were drafted concurrently in separate sessions, both claimed 0035, and the later pair were renumbered to 0037 and 0038 to preserve the already-published 0035. The gap is a renumbering artifact rather than a missing record, and it is reported rather than quietly closed because the register's value depends on it being audited, not assumed.

Resolved this period

Record	Question settled	Consequence
0035	What does the lake represent, as distinct from what it measures?	The lake is the company as a whole. Supplies the missing reasoning behind two rules already in force.
0037	May the water carry real colour, against the measured saturation ceiling?	Yes, for the water family only, with luminances pinned and bounds enforced by test. D15 stays open.
0038	May the live canvas carry scenery, and under what constraints?	Yes. Usability and per-search consistency outrank aesthetics; scenery may never look like data; hills derive from company identity alone.

Open and unanswered

Nine numbered decisions remain open on the invariants table, unchanged in count from the first issue. Three are the product owner's critical path and are carried into section 5. The remainder are not blocking today's work.

Ref	Open question	What it holds up
D9	Growth-to-speed mapping bounds	Flow-speed encoding. Excluded from the current milestone.
D10	Whether the sector rule is the right proxy for scope	Coverage edge cases and how the limit is communicated.
D12	Which period loads by default	Initial state and routing. Deferred by design; the address deliberately encodes no period.
D15	Which colour set, once colour becomes an encoding	Any use of colour as an encoding. Not needed for the current milestone.
D17	A third provenance state for filer-allocated figures	The analyst detail panel. The most direct blocker on that surface.
D18	Profit-side reconciliation rule	The second company. Microsoft reconciles exactly; others do not.
Q2	Trunk behaviour when the residual is positive	Not exercised by the current milestone.
Q3	First-run experience specification	The empty landing state. Blocked by D10 and D12.
Q4	Tax shown net or decomposed	Annotation layer on constrictions.

Two further records are ruled but unbuilt: 0033's orientation half — turning the whole picture vertical so it reads in one glance — and 0034, the clickable correction detail, whose data is already shipped and whose remaining work is interface only. 0032, the money-losing segment, is pending the product owner rather than the team.

4. Risk Register

Risks carried from the first issue are restated with any change in impact or likelihood. Four are new this period.

Risk	Movement	Assessment
The performance gate cannot produce a clean run. Two of nine tests fail: the uncapped twelve-lane load exceeds the pacing limit by 0.1 ms, and the regression gate measures a 33.4 ms worst frame against an allowed 31.4 ms. The 30 fps scenario additionally flags two checks without failing its test.	NEW · HIGH	The shipping capped path passes, and a bisect recorded in STATUS.md found the world-disabled control failing pacing worse than the world build, which points away from the new feature. Cheap to resolve: one clean run on an idle machine.
The re-recorded baseline may itself be contaminated. The reference recorded on 2026-08-30 captured a 25.10 ms worst frame, which is already over the 25.0 ms absolute limit.	NEW · MEDIUM	A baseline that enshrines a borderline number will keep producing ambiguous verdicts. It should be discarded and re-recorded rather than compared against.
The marketing site exists in no repository. Zero of its files are tracked. The home directory's ignore rule excludes the folder outright, so a disk failure loses the film, the photography and the build entirely.	NEW · HIGH	This is the same exposure the first report raised about the application, one week after that instance was closed. The remedy is identical and takes minutes.
Meta does not render. Its data layer is correct; the drawing layer refuses pending 0032.	Carried · unchanged	Second consecutive report. One of six filers with correct data cannot be shown.
The QA workstream has never started. Workstream 7 remains not started while the surface under test grew by 2,047 lines this period.	Carried · worsening	The gap widens with every shipped feature. The colocated suite is strong, but no independent adversarial pass has run against the world dressing.
The captured fixture corpus is stale. 15 of 19 envelopes no longer match pipeline output, and the corpus README misassigns the amendment case.	Carried · unchanged	Recorded as an open hand-off. Blocks meaningful regression testing.
The label placement solver is unbuilt.	Carried · unchanged	Named in STATUS.md as the most visible flaw on screen. Text can still collide.
Copy voice is unseeded. Every copy slot in the design system is open.	Carried · unchanged	No agent may author wording, so the interface ships correct but unlabelled in places.
The decision register has a numbering hole. 0036 does not exist.	NEW · LOW	Hygiene only. Worth one line in the register itself so a future reader does not hunt for a missing file.

5. Blockers and Decisions Required From the Product Owner

Four items need the product owner. Three of them were raised in the first issue and remain unanswered; each is marked with the number of consecutive reports it has been outstanding. This report's standing rule escalates any decision unanswered across three consecutive issues into the executive summary. None has reached that threshold; the first will next period if it stays open.

Decision needed	Recommended next step	Outstanding	Suggested by
How should a money-losing business segment be drawn inside an otherwise profitable company? Three options are written up with the misreadings each invites.	Choose one of the three, or ask for a fourth. Meta unblocks the moment this is answered.	2 reports	This week
Should the tax bill and other non-operating items appear as one combined number, or broken out?	Low urgency in isolation, but it now gates the annotation layer on every constriction.	2 reports	Before company two
Should companies outside today's technology-sector rule eventually be supported, and which?	Affects long-term scope rather than this week's work.	2 reports	No firm deadline
Does the world dressing look right? Six screenshots covering five filers plus Meta's refusal state are waiting in reports/world/.	Look at the set and either approve it or name what to change. The build is complete and unapproved until you do.	First raised	This week

6. Recommendation for the Coming Period

Four things, in order. First, put the marketing site under version control. It is the single largest unprotected asset the project holds, the fix takes minutes, and the identical risk was raised and closed on the application side one week ago.

Second, re-run the performance gate on an idle machine with no test suite and no local servers competing for the processor, discard the 2026-08-30 baseline and record a fresh one from that clean run. Until that happens the project cannot tell a real regression from a noisy measurement, and the two failing checks in this report cannot be closed either way. This is the highest-value hour available and it requires no decisions.

Third, answer the money-losing segment question. It is the oldest open item, it holds a filer whose data is already correct, and it reaches this report's escalation threshold next period.

Fourth, start the QA workstream. Workstream 7 has never begun while the codebase has grown past a thousand tests and gained an entire scenery layer. The colocated suites are written by the agents that wrote the code they test; an adversarial pass by a role that did not write it is the check that has been deferred longest.

Sources for this report: STREAMLINE-INVARIANTS.md, AGENT-PROTOCOL.md, STATUS.md, all 37 records under docs/decisions/, docs/product/, the git log for 2026-08-23 to 2026-09-01, a full run of the test suite and the Playwright performance gate on 2026-09-01, the pull request record, and the previous issue dated 2026-08-23. Where a claim could not be traced, it was reported as a gap rather than stated.